

Project Lab 2

Midterm Report

Dirk Thieme
R11636727
Texas Tech University

October 2022

Abstract

This paper describes the current progress of the Automated Recorder Project. While the requirements of the project will be discussed more in further sections, the general goal is to automate the playing of a recorder using a Raspberry Pi, which will be able to play any song as given to it through a flash drive. This project is being developed by Dirk Thieme, Mason Hadley, Jake Kiedasch, and Mariano Arce.

Table of Contents

1	Introduction	4
2	Body	6
2.1	Software	6
2.1.1	Flow of Control Chart	6
2.1.2	MP3 Ingestion	7
2.1.3	WAV Parsing and Processing	8
2.1.4	The Fast Fourier Transform	10
2.1.5	FFT Post-Processing	12
2.2	Hardware	14
2.2.1	The Recorder	14
2.2.2	Airflow	15
2.2.3	Solenoids and Servos	16
2.2.4	Power Supply and Driver PCB	18
2.2.5	Computational Power	19
3	Conclusion	20
4	Appendix	21
4.1	GANTT Chart and Budget	21
4.2	Safety and Ethics	22

List of Figures

I	Main Software Flowchart	6
II	ffmpeg Command and Output Implementation	7
III	WAV Processing Code	9
IV	An Example Waveform of a Song	10
V	Function to Calculate if Number is a Power of 2	11
VI	Function to Find the Next Highest Power of Two	12
VII	Example of the FFT Output	13
VIII	The Yamaha YRS-24B Recorder	15
IX	Compound Fan Setup	16
X	3D Recorder Model	17
XI	Solenoid Circuit Diagram	18
XII	Main PCB	19
XIII	The Raspberry Pi Model 4B	20
XIV	Current Budget	21
XV	Current GANTT Chart Progress	22

1 Introduction

The original premise of the Automated Recorder Project is to play a grade school recorder using a Raspberry Pi. The recorder should be able to play eight notes, and would be able to play a select number of songs chosen by the instructor of the course and a local area elementary school teacher. The recorder apparatus should be able to play in front of an elementary school class.

While the premise of the project was enough of a challenge, this project aims to go above and beyond the original scope by allowing the recorder to play all 27 notes within its capability and by allowing the user to dynamically chose the desired song to play with an MP3 file. To allow the recorder to play all notes, certain holes must only be covered halfway, a technique known as leaking and pinching [1]. To be able to play any MP3 file provided, the software will be able to process the song on demand and convert the sound wave into a collection of frequencies which correspond to notes. These notes will be automatically converted into notes that the recorder can play.

The brain of the apparatus is the Raspberry Pi 4B (furthermore Pi), a single-board computer which executes the code and controls the hardware. The Pi is powered using a quad core 64-bit ARM-Cortex A72 running at 1.5GHz [2] which brings enough power to the table in order to complete the challenging digital signal processing required. The GPIO pins are also utilized to be able to control the hardware which actuates the recorder.

The recorder is physically controlled through the use of six solenoids, three servo motors, and three compounded fans. Five of the solenoids cover and uncover the main holes, and three of the servos perform the pinching of the holes which require multiple positions. The

one other solenoid opens and closes the main labium, which stops the recorder from playing to play rests. The solenoids require 12V power, while the servos and Pi require 5V power. The Pi only outputs 5V at 16mA on the GPIO pins [2], therefore a circuit has been designed in order to be able to actuate the solenoids while not burning out the GPIO pins. To power the entire setup, a 12V, 150W power supply will be able to accept AC wall power, which it can then send to a custom made power distribution PCB, which incorporates a 12V-5V buck converter and the various high power outputs, along with the solenoid circuits.

While the structure of the software is rather simple, the implementation is exceedingly complex, and requires immense amounts of research and assistance to complete. The software is not currently finalized, however in its current state it ingests an MP3 file then converts it into a WAV file. The WAV file bytestream is parsed into its proper values, then its time-domain signal is broken up into bins which are individually processed using a Fast Fourier Transform into their frequency-domain counterpart. The magnitudes of each bin are then calculated, which reveal the peak frequency of the bin, corresponding to the note at that time period. These notes are then converted into their recorder playable counterparts, which will then be played for the correct amount of time.

2 Body

The current state of this project is comprised of the following components, the Software, a Python codebase which consists of the MP3 ingestion, the WAV parsing, the Fast Fourier Transform, and the musical outputs, and the Hardware, which consists of the recorder itself, the control devices, the air production, and the Pi which runs the whole thing.

2.1 Software

2.1.1 Flow of Control Chart

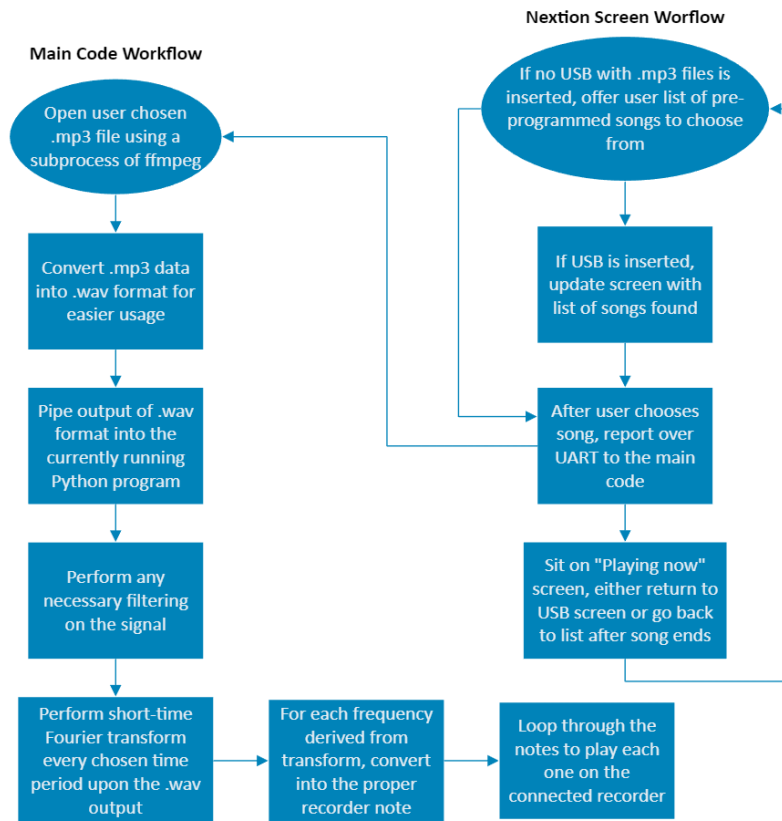


Figure I: Main Software Flowchart

As shown in Figure I, this flowchart is the initial path which the code takes. The Nextion

workflow has not yet been implemented, as it is unnecessary in determining if the main important code is functional, which itself cannot be yet determined without the construction of the apparatus. As stated previously, the code currently begins at the top if the "Main Code Workflow" section, however in the future it will begin at the "Nextion Screen Workflow". The Nextion screen is a touch screen with a built in microcontroller, which will be used to take user input and send it over to the Pi for actual use.

2.1.2 MP3 Ingestion

The bulk of the MP3 ingestion is performed by the software **ffmpeg**, originally developed by Fabrice Bellard and Bobby Bingham on December 20, 2000 [3]. Shown in Figure II, certain

```
# MAIN PROGRAM
def main():
    ... # define command to run when decoding mp3
    ... cmd = [exe, ... # ffmpeg executable
    ...       "-i", ... # input
    ...       filename, ... # name of input file (mp3)
    ...       "-ab", ... # set bitrate
    ...       "128k", ... # bitrate of input file
    ...       "-acodec", ... # set audio codec
    ...       "pcm_s16le", ... # using PCM signed 16-bit little-endian
    ...       "-map", ... # map audio streams
    ...       "0:a", ... # from file 0, take only audio
    ...       "-sn", ... # ignore any subtitle
    ...       "-vn", ... # ignore any video
    ...       "-ac", ... # set channels
    ...       "1", ... # mono
    ...       "-y", ... # override output
    ...       "-f", ... # set output format
    ...       "wav", ... # wav format
    ...       "pipe:1"] ... # pipe output to stdout and stderr

    ... # subprocess command, piping outputs to this program
    ... output = subprocess.run(cmd, stdout=subprocess.PIPE,
    ...                          stderr=subprocess.PIPE, bufsize=10**8)
```

Figure II: ffmpeg Command and Output Implementation

command flags must be used upon executing the program in order to ensure proper output

[3]. Some of the commands are less important, such as "-sn", which ignores subtitles or "-vn", which ignores any video streams (although this does implicate the software could feasibly process video data). Other commands are a necessity for proper output, such as "-f" and "wav", which sets the output format to a WAV file, and "-acodec" and "pcm_s16le", which sets the output codec correctly. The most important flag there is "pipe:1" which outputs the WAV bytestream to the stdout, instead of dumping the output to a file. All of the output is contained in the variable conveniently named "output." Output consists of two properties, stderr and stdout. As stated above, stdout contains the bytestream of the output, but for debugging purposes stderr is useful as it contains the output of the command.

2.1.3 WAV Parsing and Processing

Now that a WAV "file" has been obtained, work must be done in order to ascertain the correct values. The codec which was used to process the song is known as PCM, and the actual values that were output were unsigned 16-bit little-endian chunks. If one simply tried to read the bytes as they were, nothing but gibberish would be produced. This is because the bytes are little-endian, meaning the lower 8-bits are placed ahead of the upper 8-bits, effectively reversing the number in memory. As an example, take the decimal number 2,500. Its representation in a 16-bit hexadecimal is $0x09C4$. Storing this value into contiguous memory locations would result in the following:

$$0xC4 \rightarrow 0x09 \rightarrow \dots$$

Attempting to get this value by only reading from left to right would return the number `0xC409`, which is the decimal number 50,185, the incorrect value.

To combat this, the Python package `struct` is used, as the programmer can tell the `struct` unpacking object what format to retrieve the bytestream as. This is done using format strings, in this case the string `"<L"`. The `<` symbol tells the object the bytes will be little-endian, and the `L` tells the object to pack them into a 4-byte number [4]. This initial reading of the file allows the program to find the necessary sections of the file to read from.

WAV files consist of a standardized structure [5]. They begin with a 44-76 byte header, depending on the WAV file creator. The program parses this header to either find the "fmt" section or the "data" section. The "fmt" section contains useful file information, such as number of channels, sample rate, and codec type. The "data" section contains, quite simply, all the actual song data. The program parses the "fmt" section for the information it needs to later process the "data" section. This is shown in below in Figure III. After this is

```
def _unpackWAVE(rawSignal):
    # convert bytes to file object
    wavefile = io.BytesIO(rawSignal)

    # ignore first 4 bytes of header (RIFF header, will always be)
    _ = wavefile.read(4)
    lengthOfFile = struct.unpack(Signal.FMT_4B_LE, wavefile.read(4))[0] # get length
    # ignore next 4 bytes (denotes if WAVE file, will always be)
    _ = wavefile.read(4)

    # read file until end of file, plus 8 because we read the first 8 bytes already
    while wavefile.tell() < lengthOfFile + 8:
        # find parent of the current block
        parent = wavefile.read(4)
        # find length of the remainder of block
        try:
            currentLength = struct.unpack(Signal.FMT_4B_LE, wavefile.read(4))[0]
        except struct.error:
            # if struct no longer able to read, most likely out of bounds
            break

        # decide what to do with the current block
        if parent == b'fmt ':
            fmtData = wavefile.read(currentLength)
            # parse fmtData
            fmtTag, numChannels, framerate, avgBytesPerSecond, blockAlignVal, bitsPerSample = struct.unpack(
                Signal.FMT_FMT, fmtData)
        elif parent == b'data':
            # grab raw signal data
            readSignal = wavefile.read(currentLength)
            parsedSignal = Signal._unpackSignal(readSignal, blockAlignVal)
        else:
            # if unneeded info, skip current block
            wavefile.seek(currentLength, 1)

    return parsedSignal
```

Figure III: WAV Processing Code

complete, the variable "parsedSignal" is returned to the main program, which contains the actual waveform, and is then ready for further processing.

2.1.4 The Fast Fourier Transform

While the program now has a long list of numbers which correspond to the waveform, shown below, the waveform itself is not a very useful signal for note processing. These




Figure IV: An Example Waveform of a Song

numbers are just the amplitude of the sound at that point in time, but to find the note at that point in time, the Fourier transform of the signal must be taken. The algorithm that has been implemented for live music processing is the Cooley-Tukey Fast Fourier Transform. This algorithm was popularized in 1965 by James Cooley of IBM and John Tukey of Princeton, although the original implementation was found to have been discovered in 1805 by Carl Friedrich Gauss [6]. The main premise is to recursively divide the Fourier transform into smaller and smaller chunks, which are then added back together at the end to produce the

output [7]. The standard Discrete Fourier Transform is defined by the equation [6]:

$$X_k = \sum_{n=0}^{N-1} x_n e^{-j2\pi kn/N} \quad k = 0 \dots N - 1 \quad (1)$$

which is an $O(n^2)$ operation. The radix-2 algorithm brings this down to only $O(n \log n)$ through the recursion and division of the problem, leading to massive time and memory savings down the line.

A prerequisite of this operation is that the provided signal must be a power of 2, which is where the time savings come from [6]. In order to this, the length of the signal is checked using the code in Figure V, which compares the bits of the number and the number minus one. If the number is a power of 2, the result of ANDing them together would be 0, otherwise it would be any other number. Because the sample rate of the songs is 44.1kHz, and the

```
@staticmethod
def _isPowerOfTwo(n):
    return (n != 0) and (n & (n - 1) == 0)
```

Figure V: Function to Calculate if Number is a Power of 2

windows of time are 100ms long, the number is not a power of 2, therefore the next highest power of 2 must be found and the signal must be extended to the new number.

To find the next power of 2, a function which forces all the bits of the number to roll over until they are all 1's, then adds one to the number is used to calculate the next number to pad to. This is shown in Figure VI. Once the signal is appropriately padded, the actual math can begin. The main algorithm consists of five steps:

```

@staticmethod
def _calcNextPowerOf2(length):
    # decrement by one in case of already power of 2
    length -= 1

    # right shift all bits to force number to roll over
    length |= length >> 1
    length |= length >> 2
    length |= length >> 4
    length |= length >> 8
    length |= length >> 16

    # return rolled over number
    return length + 1

```

Figure VI: Function to Find the Next Highest Power of Two

1. Find the length of the current input list. If the length is one, return the list as is, otherwise continue.
2. Using the length of the list, calculate ω using the formula $\omega = e^{j2\pi/N}$ where N is the length.
3. Break the current list into its odd and even elements.
4. Pass the two new lists recursively into function.
5. When the lists return, recombine the odd and even lists into a single list by raising ω to the current index of the loop and doing some other math.

After these steps are completed the Fast Fourier Transform of the current list is complete.

2.1.5 FFT Post-Processing

Once the waveform is converted from a time domain signal into a frequency domain signal, further work must be done to transform the current values into usable frequencies.

To do this, the highest magnitudes of the bins of values must be calculated. First, the second half of the signal is discarded, as the second half is just the complex conjugate of the first half [6]. Then, the absolute value at each point is taken, as this is a real signal, the imaginary and negative components are irrelevant. Finally, the actual magnitude is calculated. To do this, loop through the remainder of the list, and at each even and odd point treat the even as the real component and the odd as the imaginary, then combine them using the following formula:

$$Magnitude = \sqrt{Re^2 + Im^2} \quad (2)$$

These magnitudes are then inserted into their own list, where the largest value corresponds to the highest peak in the frequency spectrum, which is almost always the frequency at that point in time. An example of the peaks is shown below in Figure VII. To find the actual




Figure VII: Example of the FFT Output

frequency associated with the peak, the equation as shown below can be used:

$$Frequency = \frac{(index\ of\ largest\ peak) * (sample\ rate)}{(length\ of\ bin)} \quad (3)$$

To actually play a song, the frequencies are placed into a list, where repeated frequencies are counted and divided by 10 to calculate the time required to play that note. To find the note, a ± 5 percent range is calculated around the frequency, which is used as a key for a dictionary with the frequencies to notes built in.

Once those notes are found, those notes are rounded up and down into the actual playable notes on the recorder. This is done by comparing the note to a list of notes which the recorder can physically play.

2.2 Hardware

2.2.1 The Recorder

As the main goal of the project is to play a recorder automatically, the most important facet of the hardware is the recorder itself. The YRS-24B is the instrument used. This recorder is fairly standard, and is tuned in C. Figure VIII depicts the recorder. This specific recorder can play notes from C5 to D7, which is a range of 27 total notes. Playing is accomplished by blowing into the mouthpiece and covering and uncovering the holes as required by each note and octave. This recorder also has a rotating end piece, which makes it easier to articulate with rigid mechanical components.



Figure VIII: The Yamaha YRS-24B Recorder

2.2.2 Airflow

The recorder is a woodwind instrument [1], and being so it requires airflow to play. The proper airflow requirement is around 0.1 to 1 liter per second [8], therefore while the speed of the air is less important, generating the correct airflow is necessary for good sound. To produce a balance between too much and too little airflow, three blower fans are mounted in series in order to control the flow. This is shown in Figure IX. The larger fan is a Dayton 1TDN1A, a 115V AC fan that can output 13CFM [9]. This blower motor provides the majority of the airflow, which is then pressurized using two 75x30mm blower fans, whose job it is only force the air through at a higher pressure.

Controlling the air is done through a solid-state relay and the Raspberry Pi, as exper-

open states, the three servos are positioned under the recorder to close and open them as needed. A 3D model of the system is shown below in Figure X.




Figure X: 3D Recorder Model

The solenoids are rated for 12V, 1A at around 500g of actuation force. In reality, that is closer to around 600mA, but everything designed is rated for 1A. As stated previously, the Pi can only output 5V at 16mA, so a circuit has been designed in order to drive them. The power supply, which will be discussed later, provides both 12V and 5V, so a MOSFET is activated with the 5V GPIO signal, which allows 12V power to run the solenoids. The solenoids are inductive loads, so when they are deactivated, a spike of current flows back into the circuit. To protect the main circuit, a flyback diode is used to stop the current. This is shown in Figure XI. For the flyback diode, a Schottky diode is used for its fast switching




Figure XI: Solenoid Circuit Diagram

abilities and high reverse breakdown voltage. When the GPIO pin is high, 12V flows through the circuit, which then can enter the solenoid. The servos are connected to dedicated 5V outputs, and the positions are controlled with the dedicated PWM connection.

2.2.4 Power Supply and Driver PCB

The PCB has been designed to provide power and outputs for all connected hardware. It is capable of providing a maximum of 12V at 7A and 5V at 6A. Input power is provided from a 12V, 15A AC to DC power supply, and is connected to the board using a GX-16 6 pin connector. A switch connects this input to the rest of the board, which then sends the power to the rest of the board.




Figure XII: Main PCB

The 5V is generated with a TPS566247DRLR, a 5V, 6A switch mode power supply. This voltage is output to four connectors, three for the servos and one for the Pi. The rest of the raw 12V power is routed to the six solenoid drivers and two fan drivers, which is the same circuit as the solenoid driver but lacking a flyback diode. The board is a four layer design, with a stackup consisting of signal-ground-12V-signal, measuring 1.6mm thick. The board is 100x60mm, with four M3 bolt holes for mounting to the apparatus.

2.2.5 Computational Power

The Pi used is a 4GB model with a quad core ARM processor. To aid the speed of the FFT execution for songs, the process will be multithreaded using Python's builtin threading library. GPIO pins 4, 17, 27, 22, 5, 6, 13, 19 are connected to the various solenoid and servo outputs. To provide the least overhead possible, instead of installing a full-fledged operating system, a headless version of Raspbian is used in combination with SSH for access.




Figure XIII: The Raspberry Pi Model 4B

3 Conclusion

As of the due date of this paper, the project is currently unfinished. That being the case, the most challenging facets, the software and airflow, have been taken care of and are ready for integration with the remainder of the parts. Some of the main components still needed are the PCB, which also needs to be assembled, a final setup of the Raspberry Pi, full assembly of the entire system, and some builtin songs in the case of no user provided songs. The code is also currently written for only MP3s, so it must be adapted to either user MP3s or to play the hard-coded songs. In addition, a touchscreen must still be implemented for user interfacing, along with a USB insertion detection script. As this project has remained fairly ahead of schedule, considerations are also being made for a MIDI interface, which can plug directly into the Pi and be used to play music on the fly, for example using a keyboard and playing a recorder note. There are also plans to get the solenoid connectors laser cut out of aluminum, and to 3D print the necessary enclosures and holders.

4 Appendix

4.1 GANTT Chart and Budget

Project Lab 1	Running Total			Total Estimate		
Direct Labor:						
<i>Category / Individual</i>	<i>Rate/Hr</i>	<i>Hrs</i>	<i>Indv. Total</i>	<i>Rate/Hr</i>	<i>Hrs</i>	<i>Indv. Total</i>
Mason	\$18.00	74	\$1,332.00	\$18.00	130	\$2,340.00
Dirk	\$18.00	74	\$1,332.00	\$18.00	130	\$2,340.00
Jake	\$18.00	72	\$1,296.00	\$18.00	130	\$2,340.00
Mariano	\$18.00	64	\$1,152.00	\$18.00	130	\$2,340.00
DL Subtotal (DL)		<i>Subtotal:</i>	\$5,112.00		<i>Subtotal:</i>	\$9,360.00
Labor Overhead	<i>rate:</i>	100%	\$5,112.00	<i>rate:</i>	100%	\$9,360.00
Total Direct Labor (TDL)			\$10,224.00			\$18,720.00
Contract Labor:						
Lab Assist	\$40.00	0	\$0.00	\$40.00	3	\$120.00
Misc.	\$18.00	0	\$0.00	\$18.00	5	\$90.00
Instructor	\$200.00	2	\$400.00	\$200.00	3	\$600.00
Total Contract Labor (TCL)			\$400.00			\$810.00
Direct Material Costs:						
Total DMC (TDM)			\$167.15			\$500.00
Equipment Rental Costs:	<i>Value:</i>	<i>Rental Rate:</i>		<i>Value:</i>	<i>Rental Rate:</i>	
Oscilloscope	\$1,275.00	0.20%	\$135.15	\$1,275.00	0.20%	\$234.60
Power Supply	\$966.00	0.20%	\$102.40	\$966.00	0.20%	\$177.74
Multimeter	\$777.00	0.20%	\$82.36	\$777.00	0.20%	\$142.97
Waveform Generator	\$1,051.00	0.20%	\$111.41	\$1,051.00	0.20%	\$193.38
Total Rental Costs (TRM)			\$431.31			\$748.70
Total TDL+TCL+TDM+TRM			\$11,222.46			\$20,778.70
Business Overhead		55%	\$6,172.36		55%	\$11,428.28
Total Cost:		<i>Current:</i>	\$17,394.82		<i>Estimate:</i>	\$32,206.98

Figure XIV: Current Budget

As shown in Figures XIV and XV, the budget and GANTT are progressing well. The budget is currently sitting at around 63% used, which is about how much of the semester has currently passed, keeping the budget on track. As shown in the GANTT chart, the majority of the items to be completed have either been completely finished or at least halfway done.

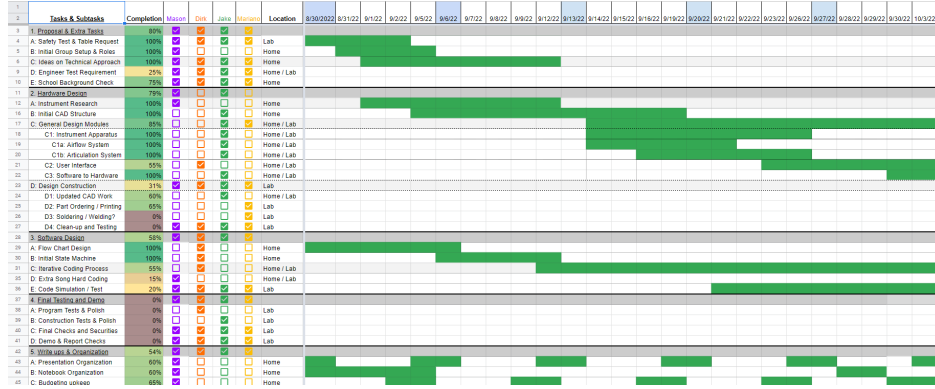


Figure XV: Current GANTT Chart Progress

4.2 Safety and Ethics

Special considerations must be thought of safety wise, as this project is meant to be presented to a classroom of elementary schoolers. Because of this, care has been taken to ensure AC power is never exposed anywhere a student might touch, and that the touchscreen access in general should only be relegated to the teacher of the classroom or the people of the lab group. It is also important to consider volume levels, as overblowing the recorder could cause a loud, sharp noise which could hurt a student. The maximum airflow of the fans has been found, and it is within an acceptable range as to not hurt any hearing. The ethical ramifications stem from the fact that human musicians are already paid very little for their talents, so developing an automated system to play an instrument might lead to less jobs. However, this machine will never be able to be as versatile and talented as a real person playing, so I think it is safe to say this shouldn't hurt any chances.

References

- [1] “Recorder (musical instrument).” (Oct. 19, 2022), [Online]. Available: [https://en.wikipedia.org/wiki/Recorder_\(musical_instrument\)#Acoustics](https://en.wikipedia.org/wiki/Recorder_(musical_instrument)#Acoustics).
- [2] *Raspberry Pi 4 Model B*, Raspberry Pi (Trading) Ltd., Jun. 2019.
- [3] “ffmpeg Documentation.” (Oct. 19, 2022), [Online]. Available: <https://ffmpeg.org/ffmpeg.html>.
- [4] “struct — Interpret bytes as packed binary data.” (Oct. 19, 2022), [Online]. Available: <https://docs.python.org/3/library/struct.html>.
- [5] “What is a WAV file?” (Oct. 19, 2022), [Online]. Available: <https://docs.fileformat.com/audio/wav/>.
- [6] “Fast Fourier transform.” (Oct. 19, 2022), [Online]. Available: https://en.wikipedia.org/wiki/Fast_Fourier_transform.
- [7] J. W. Cooley and J. W. Tukey, “An Algorithm for the Machine Calculation of Complex Fourier Series,” 1965.
- [8] J. Wolfe. “Air speed and blowing pressure in woodwind and brass instruments: how important are they?” (Oct. 24, 2022), [Online]. Available: <https://www.phys.unsw.edu.au/jw/air-speed.html>.
- [9] “OEM Blower: Round, With Flange, 2 in Blower Wheel Dia, Shaded Pole Motor, 115V AC, Auto.” (Oct. 24, 2022), [Online]. Available: <https://www.grainger.com/product/DAYTON-OEM-Blower-Round-1TDN1>.